



LUCENE / SOLR REVOLUTION / 2016

<http://lucenerevolution.com>

A dark silhouette of a city skyline, likely Boston, featuring various skyscrapers and a church steeple, positioned along the bottom edge of the poster.

OCTOBER 11-14
BOSTON, MA

Wiring ML Models into the inner loop of a Lucene Scorer

Jake Mannix @pbrane

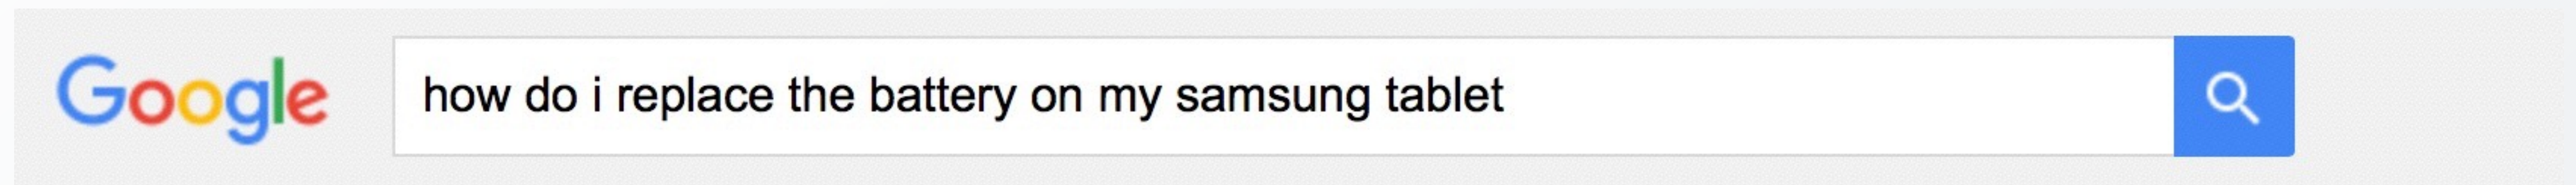
Lead Data Engineer, Lucidworks

previously: IR/RecSys @ LinkedIn,
Search & p13n @ Twitter,
Semantic Search @ Ai^2

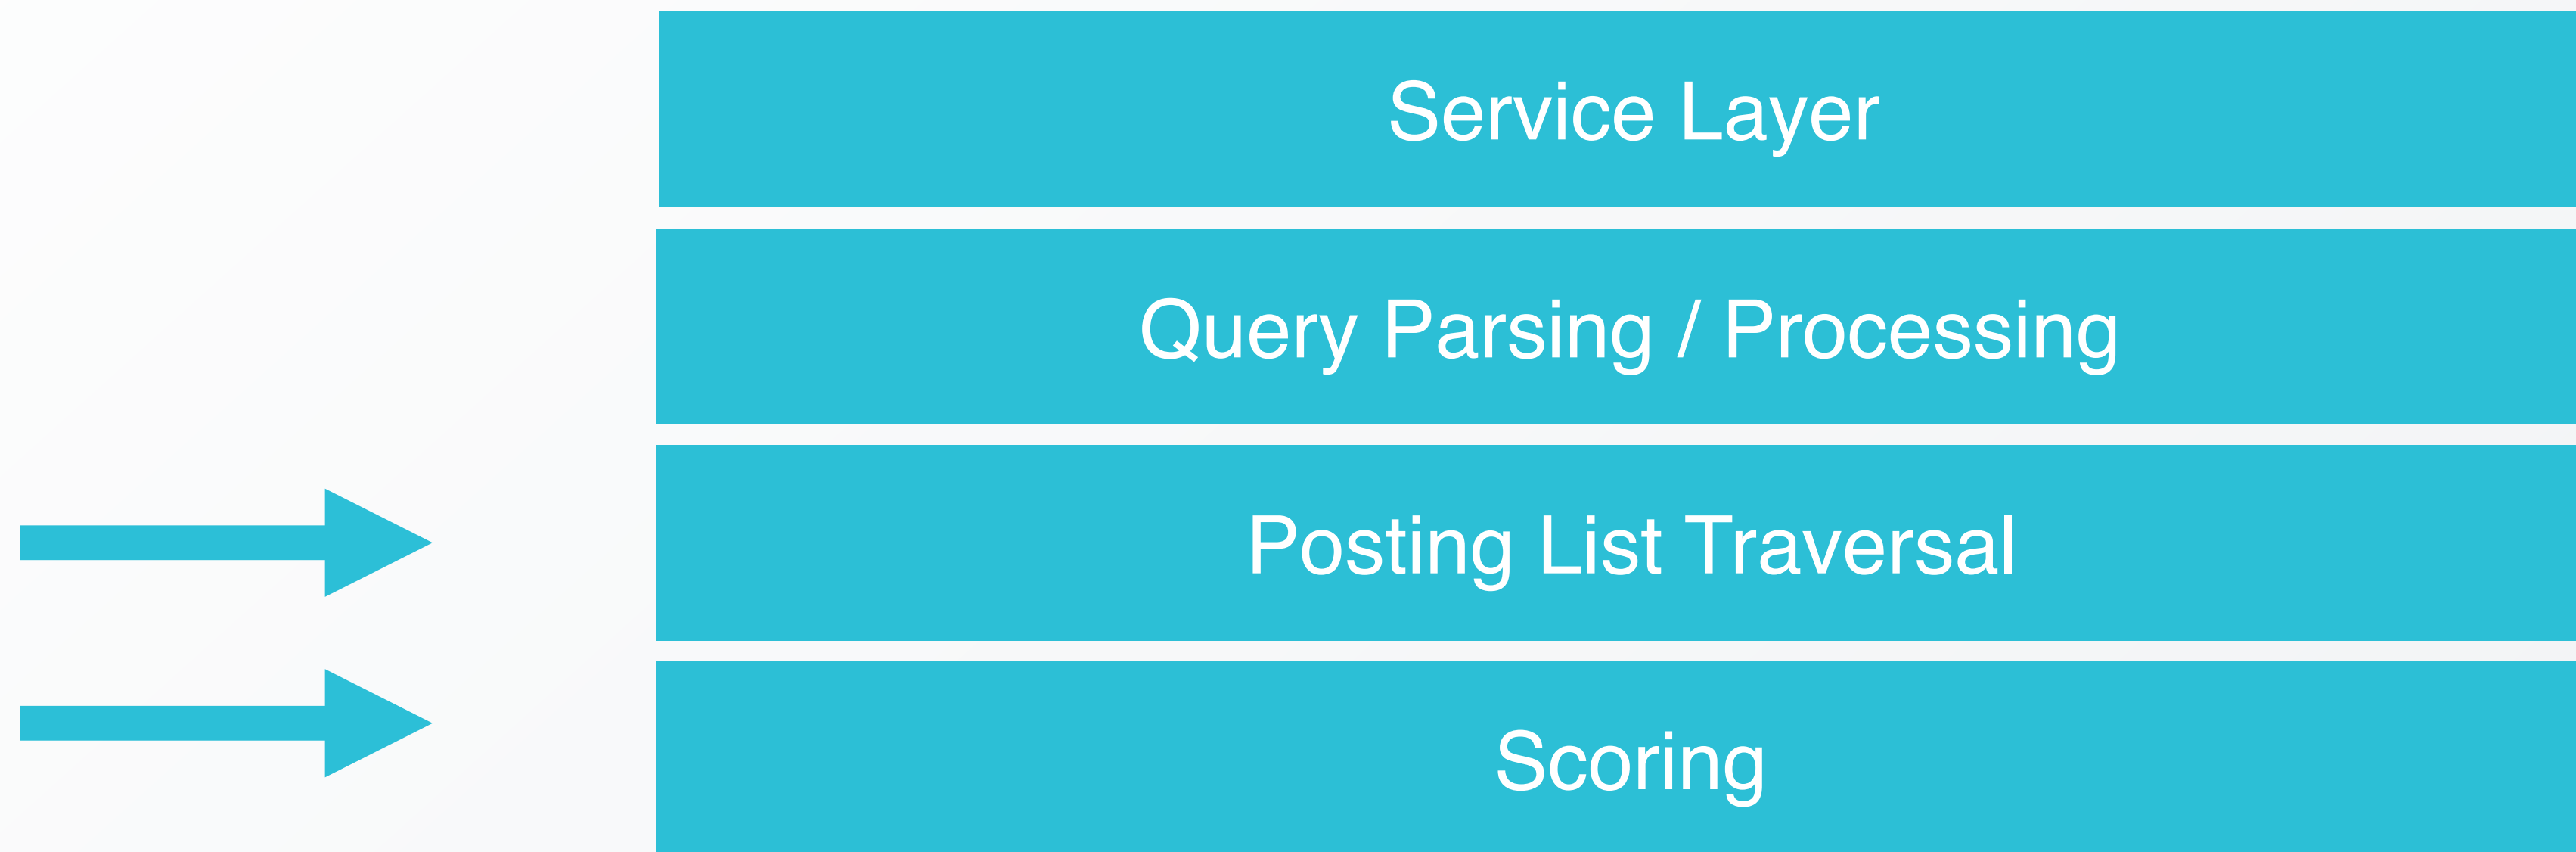
(Apache Mahout committer)

- Search: State of the Industry
- Lucene OOTB Scoring / Similarity
- Trained models for Lucene
- DisjunctionScorer
- IdentifiedDisjunctionScorer

- Google has trained users too well. Queries like these are becoming commonplace:



- NLP and query categorization / classification is often needed at query “parsing” stage
- Behind the scenes, previous queries of you (and everyone else!) factor into ranking



- Internally, queries tend to hit semi-structured indexes in increasingly complicated ways:
 - recruiter searching on LinkedIn: "sr. SDE java scala spark big data"
 - becomes:
 - `(title:("sr sde" OR "senior software development engineer" OR sde^0.1)^3.5 AND (skills:(java OR scala OR spark OR "big data")^1.5 OR skills:(hadoop OR pig OR hive)^0.1)`
 - (and that's without hitting multilingual fields, stemmed fields, cluster-label fields, synonym fields, etc...)

- Why did Google succeed (at first)?
 - PageRank - crowdsourced page relevance by the global popularity structure of the web graph
- What are dominant factors for search at Google, Amazon, Facebook, Twitter, LinkedIn?
 - popularity (both globally, by demographic, context, and socially)
 - “social proof”
 - recency (sometimes as filter, sometimes as ranking factor)
 - past click/engagement history
 - user-model / preferences
- What is not a dominant factor, beyond the basics: textual similarity

- **Query -> Weight -> Scorer**
 - Lucene **Query** is a (reusable) blueprint to tell an **IndexSearcher** how to build a **DocIdSetIterator**
 - **Weight** is an intermediate **IndexSearcher**-stateful object ready to build **Scorers**.
 - it precomputes any query-global information (query normalization, separate sub-query parts which are for scoring, and which for just matching)
 - **Scorer**: both an iterator over posting list, and computes scores for each document

- Instantiable **Query** subclasses:
 - **TermQuery**
 - **BooleanQuery**, **DisjunctionMaxQuery**
 - **WildcardQuery**, **PrefixQuery**
 - **PhraseQuery**, **MultiPhraseQuery**
 - **FuzzyQuery**, **RegexQuery**
 - **TermRangeQuery**, **PointRangeQuery**
 - **MatchAllDocsQuery**
 - Span queries: `org.apache.lucene.search.spans.*`

- Implements the generic scoring model by comparing query <-> document similarity
- **CollectionStatistics, TermStatistics**
- Helper functions (e.g. what impl for "tf()" or "idf()" do you want to use?)
- `queryNorm(float sumOfSquaredWeights)` - make scores comparable across queries
- `coord(int overlap, int maxOverlap)` - score factor from #matching terms between doc/query

- **TFIDFSimilarity**
- **BM25Similarity** (now default!)
- **LMSimilarity** (language models with different smoothing)
 - **LMDirichletSimilarity**
 - **LMJelinekMercerSimilarity**
- **DFISimilarity** (divergence from independence)
- **DFRSimilarity** (divergence from randomness)
- **IBSimilarity** (information-based models)
 - allows pluggable distributions modeling term occurrence
- combiners: **PerFieldSimilarity**, **MultiSimilarity**

- Fielded query model with trained weights applied to each field
 - $(\text{title:q_t})^{\text{title_w}}$ OR $(\text{desc:q_d})^{\text{desc_w}}$ OR $(\text{body:q_b})^{\text{body_w}}$ OR $(\text{_all:q_a})^{\text{all_w}}$
- Rerank top-K results - see Diego's talk on LTR tomorrow!
- How to score a decision tree or neural net?
 - in particular, some models want to know things like:
 - $\text{if } ((\text{titleScore} > 0.2 \ \&\& \ \text{skillsScore} > 0.1) \ \&\& \ (\text{descScore} > 0.6 \ || \ \text{bodyScore} > 0.95)) \dots$

- Decision Tree
 - CART
 - Random Forest, Bagging, Boosting, etc
 - even hand-crafted (not learned!) “artisanal” DTs (i.e. business rules)
- Neural Net
- ~~Deep Learning~~ (see line above)

- Quiz (for non-Lucene developers): How would you efficiently traverse the posting list for:
 - $q = \text{"term1"} \text{ OR } \text{"term2"} \text{ OR } \text{"term3"} \text{ OR } \dots \text{ OR } \text{"termN"} \text{ ?}$

- Answer:
 - leaf query in the boolean tree \leftrightarrow (docId-ordered) Iterator over posting list
 - put all these Iterators into a min-heap (PriorityQueue), ordered by the value of `.currentDocId()`
 - pull Iterators off of the heap, checking to see you're still on the same document, accumulate scores from each

- But what if you want to score the sub-queries other than by simply summing the results?
- Let's look at the current implementation of DisjunctionScorer in Lucene

```
package org.apache.lucene.search;

import ...

/** A Scorer for OR like queries, counterpart of ConjunctionScorer.
 */
final class DisjunctionSumScorer extends DisjunctionScorer {
    private final float[] coord;

    /** Construct a DisjunctionScorer.
     * @param weight The weight to be used.
     * @param subScorers Array of at least two scorers.
     * @param coord Table of coordination factors
     */
    DisjunctionSumScorer(Weight weight, List<Scorer> subScorers, float[] coord, boolean needsScores) {
        super(weight, subScorers, needsScores);
        this.coord = coord;
    }

    @Override
    protected float score(DisiWrapper topList) throws IOException {
        double score = 0;
        int freq = 0;
        for (DisiWrapper w = topList; w != null; w = w.next) {
            score += w.scorer.score();
            freq += 1;
        }
        return (float)score * coord[freq];
    }
}
```

- $(\text{title:q_t})^{\text{title_w}}$ OR $(\text{desc:q_d})^{\text{desc_w}}$ OR $(\text{body:q_b})^{\text{body_w}}$ OR $(\text{_all:q_a})^{\text{all_w}}$
- field boosts can encode the model weights
 - summing: linear regression scoring
 - summing then apply $\text{logit}()$: logistic regression scoring

- IdentifiedScorer*, IdentifiedDisjunctionScorer*

```
/**
 * Scorer for Disjunctions where you want to keep track of which sub-scorer
 * contributed to the current hit, and scoring can non-linearly combine once
 * all sub-scores are available
 */
abstract class IdentifiedDisjunctionScorer extends DisjunctionScorer {

    private final int numSubScorers;

    public IdentifiedDisjunctionScorer(Weight weight, List<Scorer> subScorers, boolean needsScores) {
        super(weight, subScorers, needsScores);
        numSubScorers = subScorers.size();
    }

    protected float score(IdentifiedScorer topList) throws IOException {
        double[] scores = new double[numSubScorers];
        for (IdentifiedScorer w = topList; w != null; w = w.nextIdentifiedScorer) {
            scores[topList.id] = w.scorer.score();
        }
        return combineScores(scores);
    }

    public abstract float combineScores(double[] subScorerScores);
}
```

- (do not exist yet! hack with me tomorrow?)

- Latency:
 - complex scoring at every hit - will it work on your real data?
 - (early termination of scoring function [return 0 early when you know it won't score well])
- True API compatibility
 - there are many ways **Query** and **Scorer** objects can be used (as a filter with no scoring, scoring in bulk, nesting hierarchy: **IdentifiedQuery** containing a **BooleanQuery**, or vice versa?)
- Truly comparable query norms
 - Most query scores are “kinda comparable” - but are not strictly bounded!
 - So training things like logistic regression models often get confused
- Collecting training data from implicit information (clicks and related engagement signals)
- Model I/O

- Fusion: <http://www.lucidworks.com/products/fusion>
- Search Hub: <http://searchhub.lucidworks.com>
- Company: <http://www.lucidworks.com>
 - Our blog: <http://www.lucidworks.com/blog>
- Twitter: @pbrane

Thanks!

this page left intentionally blank